



Deep Learning

From Mainstream Hype to training actual Video-Game Agents



Turn: -0.9851
Engine: 0.90426
Fitness: 0.45368

Generation: 9




Samuel Arzt

Purpose of this talk

- Understand fundamental concepts of Deep Learning / Machine Learning
- Starting point to read more about and dive deeper into the field
- Second half: Focus on Learning Algorithms applied to Video Games
- Disclaimer: Sometimes not scientifically accurate and very simplified
 - Avoided complex formulars where possible


What is Deep Learning?

- Simplified:
Combination of Deep Neural Networks and Machine Learning
or
Using Machine Learning Algorithms to train Deep Neural Networks



Deep Neural Networks

+




Machine Learning

What are Neural Networks?

<https://youtu.be/rEDzUT3ymw4>

[1]

Neural Networks




EXPLAINED IN A MINUTE

When does a NN become deep?


- Simple (and outdated) Definition:
Any Network with more than one hidden layer
- Modern Understanding:
Not well defined, „credit assignment of
paths of sufficient length“ ^[1]

„Problems of depth > 10 require
very deep learners“ ^[1]



NN Architectures

- Feedforward Neural Networks
 - Acyclic Graphs, only connections to „upper“ layers
- Recurrent Neural Networks
 - Cycles and connections to previous layers allowed
- Convolutional Neural Networks
 - Weight sharing and spatial alignment of neurons, specialized for image input; Comparable to how image-filters work
 - Many new variants since 2014



What is Machine Learning?

- Wikipedia: „Use statistical techniques to give computer systems the ability to *learn* (e.g., **progressively improve performance** on a specific task) from data, **without being explicitly programmed.**” ^[1]
- „Algorithms that can learn from and make predictions on data by making data-driven predictions or decisions, through **building a model from sample inputs**” ^[1]

What is Machine Learning?

- Traditional Algorithms:
Hard coded conditions and sequences of instructions.
- Machine Learning:
Algorithm may start with random parameters and successively learns from input data.
Goal: Improve performance (i.e. ability to solve task) over time.

Three Main Categories of ML


- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

Supervised Learning

- Learning from „labeled“ data
 - E.g.: Large dataset of already labeled cat & dog images
-> train classifier for cat / dog images
- Feed input to network -> depending on output, tweak weights to be „less wrong“ (minimize loss)
 - Stochastic gradient descent (backpropagation)
- In combination with ConvNets, most widely employed ML technique

Gradient Descent


- Gradient = direction of steepest descent for multivariate functions
 - Calculated using partial derivative of loss function
- Adjust parameters in direction of Gradient
- Stochastic GD
 - Use batch of data (random or simply order in training set) instead of entire data-set to compute gradient



Reinforcement Learning

- „Agent“ learns through interacting with Environment
 - Feedback („Reward Signal“) from Environment
 - Goal: Learn a „Policy“, which maximizes accumulated reward
 - No labels, but trial and error

- Standard RL Setting:
 - State, Action, Reward
 - Epoch / Episode
 - Continuous / Sparse / Delayed Rewards



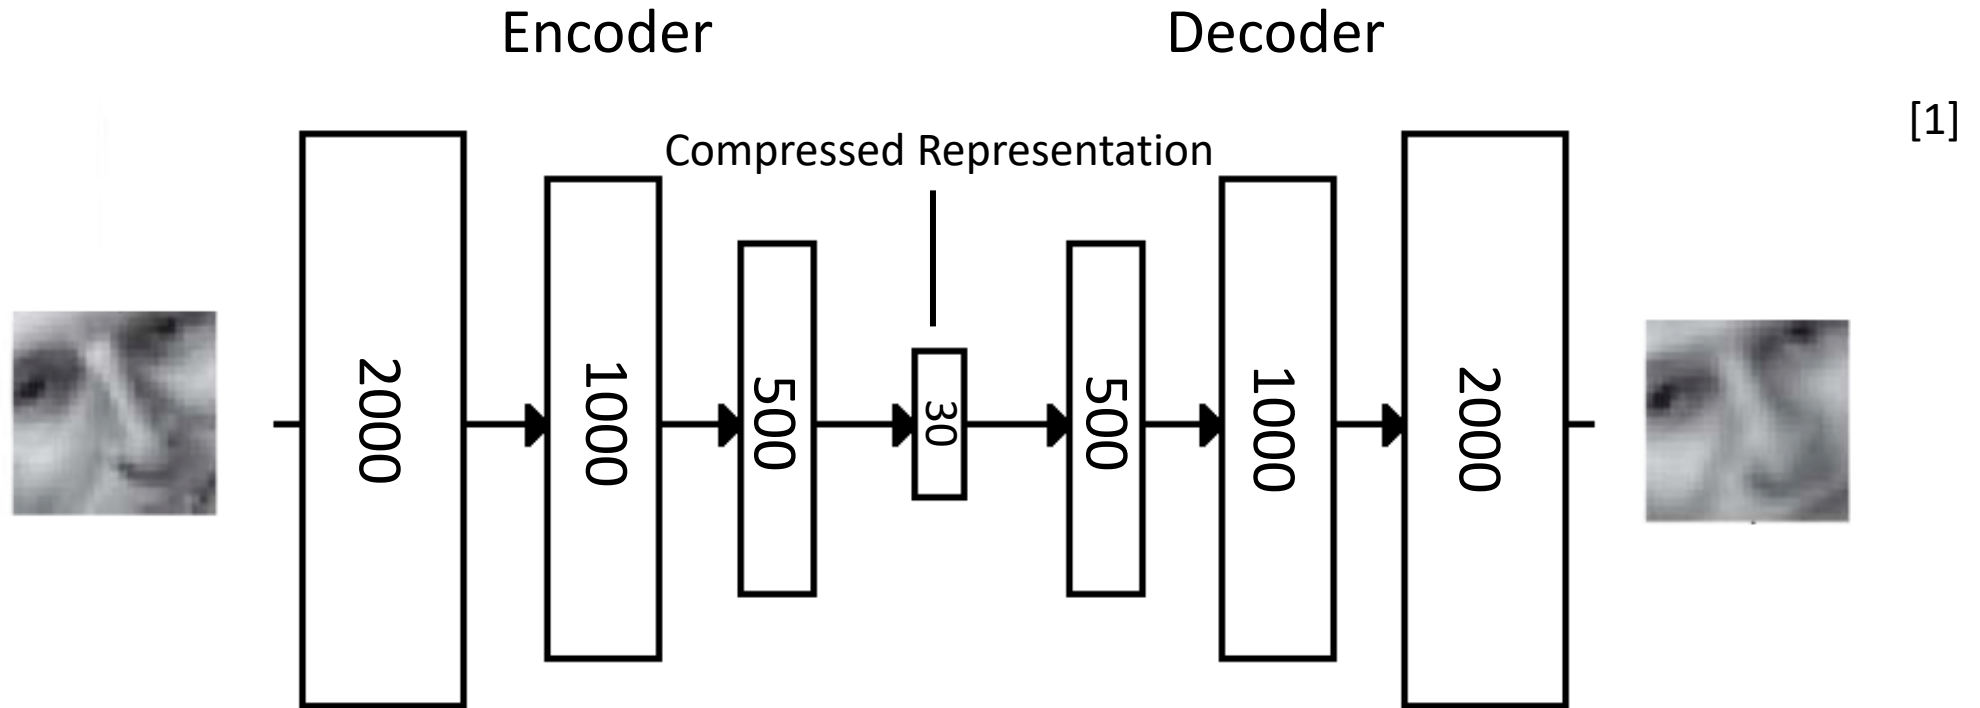
The agent–environment interaction in a Markov decision process.

- Most famous: Q-Learning

Unsupervised Learning

- Learn from completely unlabeled data
- Probably least understood / progress from all three subcategories
- Happens in Humans / Animals all the time („common sense“)
- Example:
Humans / Animals don't need to be taught Newton's 1st Law to understand and anticipate a ball falling off the table when pushed (learned solely from observation; no feedback, no labels)
- LeCun: *UL is the cake, SL the icing, RL the cherry on top.* (in respect to data efficiency)
Unfortunate situation: we know how to make the icing and cherry but not how to make the cake ^[1]

UL Example: Encoder Decoder Network

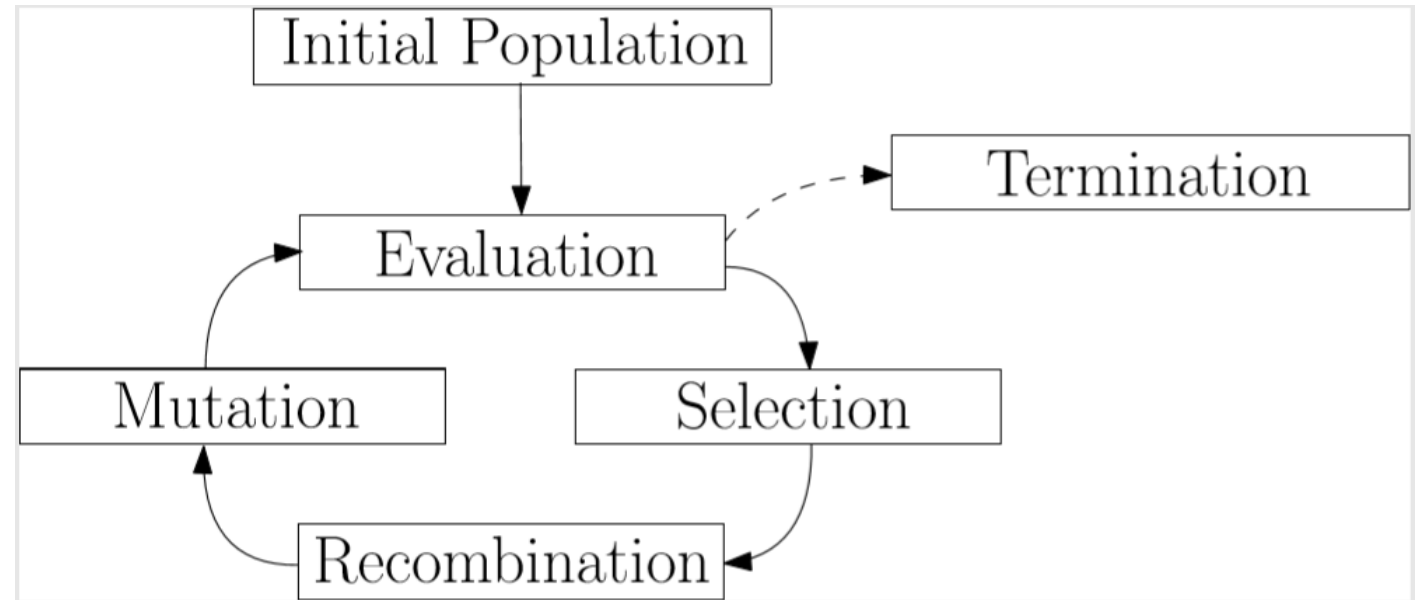


ML Categories: Summary

- Supervised Learning:
Learn using labeled data; Most popular sub-category
- Reinforcement Learning
Interact with environment and learn from feedback; Easy to integrate into typical game-setup
- Unsupervised Learning
Learn completely without labels or feedback; Currently often applied for compression / clustering

Evolutionary Algorithms

- Typical Cycle:
 - Evaluate current Population
 - Select „best“
 - Recombine selected to form new Generation
 - Mutate to introduce new information




- Examples: Genetic Algorithm, Evolution Strategies
- Which subcategory of ML is this?

Neural Networks + Evolutionary algorithms

[1]

<https://youtu.be/Aut32pR5PQA>



Q-Learning^[1]

- Traditionally most well known RL algorithm
- Action-Value Function „ Q^π “ calculates „value“ of taking a specific action , a' in state , s' under current policy , π' :

$$Q^\pi(s_t, a_t) = r_t + \gamma Q^\pi(s_{t+1}, a_{t+1})$$

- Q-Function is used to determine the „quality“ of a policy.

Q-Learning

- Q-Learning tries to approximate, i.e. „learn“, the optimal Q-Function , Q^* ‘:

$$Q^*(s_t, a_t) = r_t + \gamma \max_{a' \in \mathcal{A}} Q^*(s_{t+1}, a')$$

- Optimal policy can be derived from Q^* :

$$\pi(s_t) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s_t, a)$$

Q-Learning

- Standard Q-Learning update:

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \overbrace{\left(\underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} \right)}^{\text{learned value}}$$

[1]

- Use of ϵ -greedy policy for exploration
 - Choose optimal action with probability $(1 - \epsilon)$, otherwise random action

Deep Reinforcement Learning

- Deep Q-Network ^[1]
 - Tabular Q-Function for large state-spaces unfeasible
Idea: Use Neural Network to approximate Q-Function
- Alpha Go, Alpha Go Zero ^[2]
 - MCTS, Neural Network approximates Value function, Self-Play
- Newer Algorithms: PPO, A3C, ACKTR, Meta-Learning, ...



[1]: Mnih et. al., 2015 „Human-level control through deep reinforcement learning.“

[2]: Silver et. al., 2017 „Mastering the game of go without human knowledge.“


DQN

- Use Deep Neural Network to approximate Q^*
- Update rule similar to traditional, but now used as loss for SGD
- Use ConvNet for end-to-end learning with pixel inputs
- Requires some extra tricks to work:
 - Separate target network for loss
 - Experience Replay (sampled from for SGD batches)
- Many improvements over the past years: Double DQN, Prioritized ER, Noisy Nets, Dueling DQN, ... , Rainbow

DQN on Atari



DQN on more Complex visuals using ML-Agents ^[1]



Good starting points

- Genetic Algorithms: Darrel Whitley,
„A Genetic Algorithm Tutorial.“
- RL in general: Sutton and Barto,
„Reinforcement Learning: An Introduction“ (new version in 2017/18)
- Unity ML-Agents:
<https://unity3d.com/machine-learning>
- OpenAI Spinning Up in Deep RL (very new!):
<https://blog.openai.com/spinning-up-in-deep-rl/>



@SamuelArzt



SamuelArzt



arztsamuel.github.io