



# Deep Learning

From Mainstream Hype to training actual Video-Game Agents



Turn: -0.9851  
Engine: 0.90426  
Fitness: 0.45368

Generation: 9




Samuel Arzt

# Purpose of this talk

- Understand fundamental concepts of Deep Learning / Machine Learning
- Starting point to read more about and dive deeper into the field
- Second half: Focus on Learning Algorithms applied to Video Games
- Disclaimer: Sometimes not scientifically accurate and very simplified
  - Avoided complex formulars where possible

# What is Deep Learning?

- Simplified:  
Combination of Deep Neural Networks and Machine Learning  
or  
Using Machine Learning Algorithms to train Deep Neural Networks



Deep Neural Networks

+




Machine Learning

# What are Neural Networks?

<https://youtu.be/rEDzUT3ymw4>

[1]


## Neural Networks



**EXPLAINED IN A MINUTE**


# When does a NN become deep?

- Simple (and outdated) Definition:  
Any Network with more than one hidden layer
- Modern Understanding:  
Not well defined, „credit assignment of  
paths of sufficient length“ <sup>[1]</sup>  
  
„Problems of depth  $> 10$  require  
*very deep* learners“ <sup>[1]</sup>



# NN Architectures

- Feedforward Neural Networks
  - Acyclic Graphs, only connections to „upper“ layers
- Recurrent Neural Networks
  - Cycles and connections to previous layers allowed
- Convolutional Neural Networks
  - Weight sharing and spatial alignment of neurons, specialized for image input; Comparable to how image-filters work
  - Many new variants since 2014



# What is Machine Learning?

- Wikipedia: „Use statistical techniques to give computer systems the ability to *learn* (e.g., **progressively improve performance** on a specific task) from data, **without being explicitly programmed.**” <sup>[1]</sup>
- „Algorithms that can learn from and make predictions on data by making data-driven predictions or decisions, through **building a model from sample inputs**” <sup>[1]</sup>

# What is Machine Learning?

- Traditional Algorithms:  
Hard coded conditions and sequences of instructions.
- Machine Learning:  
Algorithm may start with random parameters and successively learns from input data.  
Goal: Improve performance (i.e. ability to solve task) over time.



# Three Main Categories of ML


- Supervised Learning
- Unsupervised Learning
- Reinforcement Learning

# Supervised Learning

- Learning from „labeled“ data
  - E.g.: Large dataset of already labeled cat & dog images  
-> train classifier for cat / dog images
- Feed input to network -> depending on output, tweak weights to be „less wrong“ (minimize loss)
  - Stochastic gradient descent (backpropagation)
- In combination with ConvNets, most widely employed ML technique

# Gradient Descent


- Gradient = direction of steepest descent for multivariate functions
  - Calculated using partial derivative of loss function
- Adjust parameters in direction of Gradient
- Stochastic GD
  - Use batch of data (random or simply order in training set) instead of entire data-set to compute gradient



# Reinforcement Learning

- „Agent“ learns through interacting with Environment
  - Feedback („Reward Signal“) from Environment
  - Goal: Learn a „Policy“, which maximizes accumulated reward
  - No labels, but trial and error

- Standard RL Setting:
  - State, Action, Reward
  - Epoch / Episode
  - Continuous / Sparse / Delayed Rewards



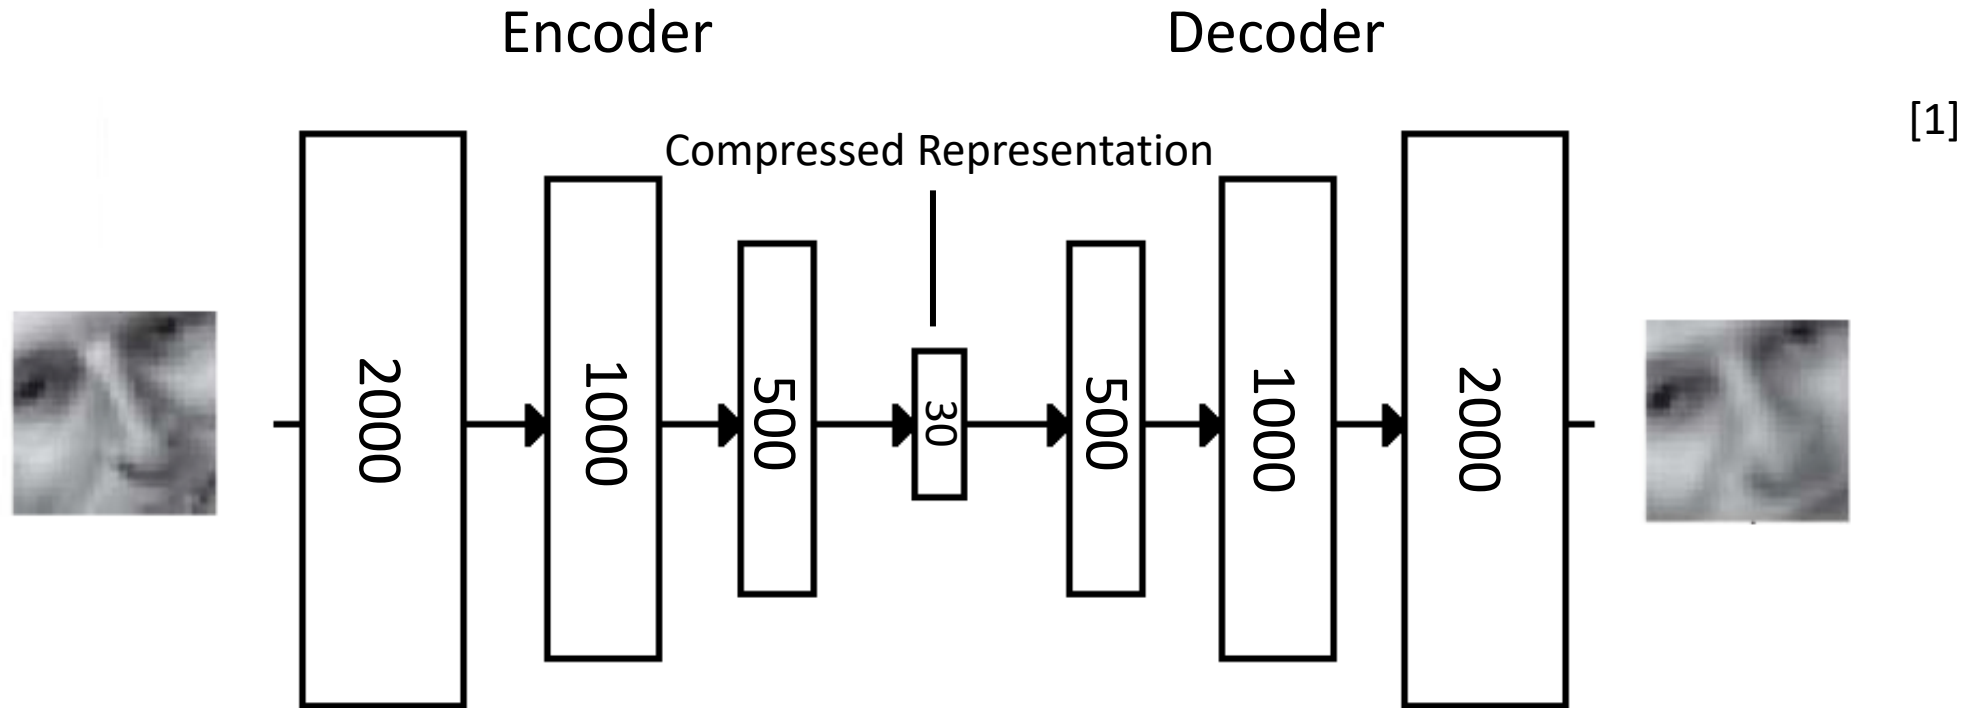
The agent–environment interaction in a Markov decision process.

- Most famous: Q-Learning

# Unsupervised Learning

- Learn from completely unlabeled data
- Probably least understood / progress from all three subcategories
- Happens in Humans / Animals all the time („common sense“)
- Example:  
Humans / Animals don't need to be taught Newton's 1st Law to understand and anticipate a ball falling off the table when pushed (learned solely from observation; no feedback, no labels)
- LeCun: *UL is the cake, SL the icing, RL the cherry on top.* (in respect to data efficiency)  
*Unfortunate situation: we know how to make the icing and cherry but not how to make the cake* <sup>[1]</sup>

# UL Example: Encoder Decoder Network

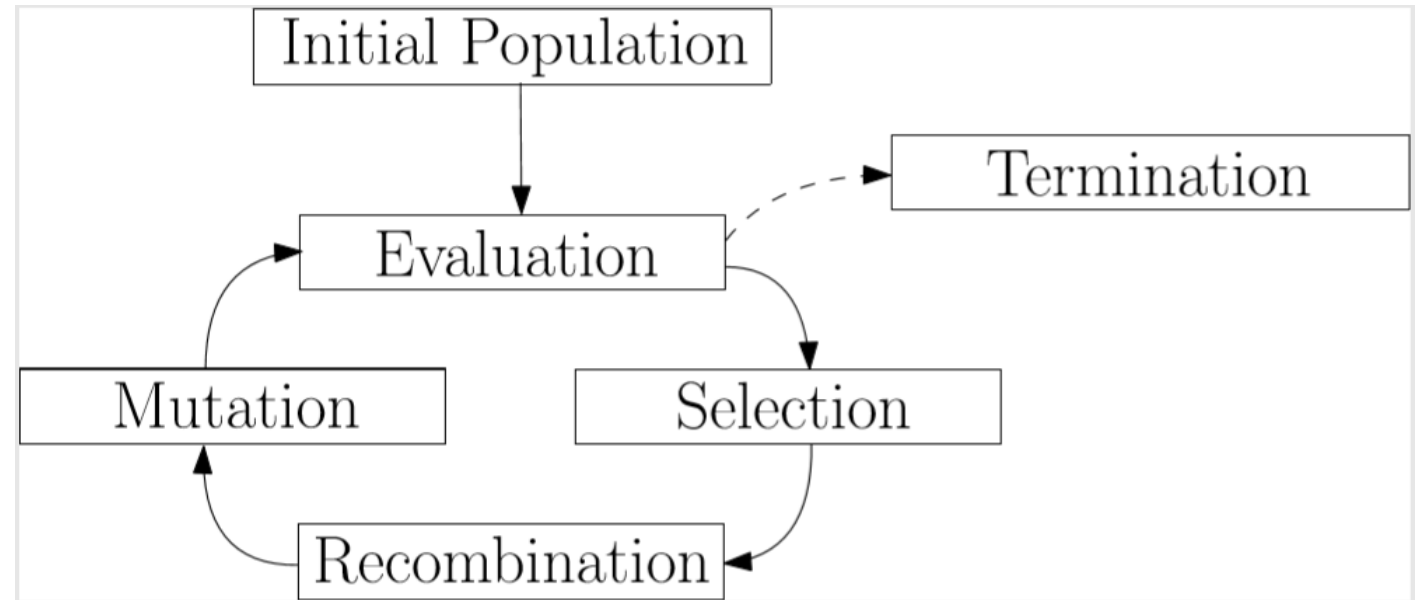


# ML Categories: Summary

- Supervised Learning:  
Learn using labeled data; Most popular sub-category
- Reinforcement Learning  
Interact with environment and learn from feedback; Easy to integrate into typical game-setup
- Unsupervised Learning  
Learn completely without labels or feedback; Currently often applied for compression / clustering

# Evolutionary Algorithms

- Typical Cycle:
  - Evaluate current Population
  - Select „best“
  - Recombine selected to form new Generation
  - Mutate to introduce new information




- Examples: Genetic Algorithm, Evolution Strategies
- Which subcategory of ML is this?



# Neural Networks + Evolutionary algorithms

[1]

<https://youtu.be/Aut32pR5PQA>



[1]: <https://arztsamuel.github.io/en/projects/unity/deepCars/deepCars.html>

# Q-Learning<sup>[1]</sup>

- Traditionally most well known RL algorithm
- Action-Value Function „ $Q^\pi$ “ calculates „value“ of taking a specific action , $a'$  in state , $s'$  under current policy , $\pi'$ :

$$Q^\pi(s_t, a_t) = r_t + \gamma Q^\pi(s_{t+1}, a_{t+1})$$

- Q-Function is used to determine the „quality“ of a policy.

# Q-Learning

- Q-Learning tries to approximate, i.e. „learn“, the optimal Q-Function  $Q^*$ :

$$Q^*(s_t, a_t) = r_t + \gamma \max_{a' \in \mathcal{A}} Q^*(s_{t+1}, a')$$

- Optimal policy can be derived from  $Q^*$ :

$$\pi(s_t) = \operatorname{argmax}_{a \in \mathcal{A}} Q^*(s_t, a)$$

# Q-Learning

- Standard Q-Learning update:

$$Q^{new}(s_t, a_t) \leftarrow (1 - \alpha) \cdot \underbrace{Q(s_t, a_t)}_{\text{old value}} + \underbrace{\alpha}_{\text{learning rate}} \cdot \overbrace{\left( \underbrace{r_t}_{\text{reward}} + \underbrace{\gamma}_{\text{discount factor}} \cdot \underbrace{\max_a Q(s_{t+1}, a)}_{\text{estimate of optimal future value}} \right)}^{\text{learned value}}$$

[1]

- Use of  $\epsilon$ -greedy policy for exploration
  - Choose optimal action with probability  $(1 - \epsilon)$ , otherwise random action

# Deep Reinforcement Learning

- Deep Q-Network <sup>[1]</sup>
  - Tabular Q-Function for large state-spaces unfeasible  
Idea: Use Neural Network to approximate Q-Function
- Alpha Go, Alpha Go Zero <sup>[2]</sup>
  - MCTS, Neural Network approximates Value function, Self-Play
- Newer Algorithms: PPO, A3C, ACKTR, Meta-Learning, ...



[1]: Mnih et. al., 2015 „Human-level control through deep reinforcement learning.“

[2]: Silver et. al., 2017 „Mastering the game of go without human knowledge.“


# DQN

- Use Deep Neural Network to approximate  $Q^*$
- Update rule similar to traditional, but now used as loss for SGD
- Use ConvNet for end-to-end learning with pixel inputs
- Requires some extra tricks to work:
  - Separate target network for loss
  - Experience Replay (sampled from for SGD batches)
- Many improvements over the past years: Double DQN, Prioritized ER, Noisy Nets, Dueling DQN, ... , Rainbow

# DQN on Atari



# DQN on more Complex visuals using ML-Agents <sup>[1]</sup>





# Good starting points

- Genetic Algorithms: Darrel Whitley,  
„A Genetic Algorithm Tutorial.“
- RL in general: Sutton and Barto,  
„Reinforcement Learning: An Introduction“ (new version in 2017/18)
- Unity ML-Agents:  
<https://unity3d.com/machine-learning>
- OpenAI Spinning Up in Deep RL (very new!):  
<https://blog.openai.com/spinning-up-in-deep-rl/>



@SamuelArzt



SamuelArzt



arztsamuel.github.io